

# Garbage collection

# GC

- Every modern programming language allows programmers to allocate new storage dynamically
  - New records, arrays, tuples, objects, closures, etc.
- Every modern language needs facilities for reclaiming and recycling the storage used by programs
- It's usually the most complex aspect of the run-time system for any modern language

# Memory layout

per process  
virtual memory

new pages allocated  
via calls to OS

physical memory

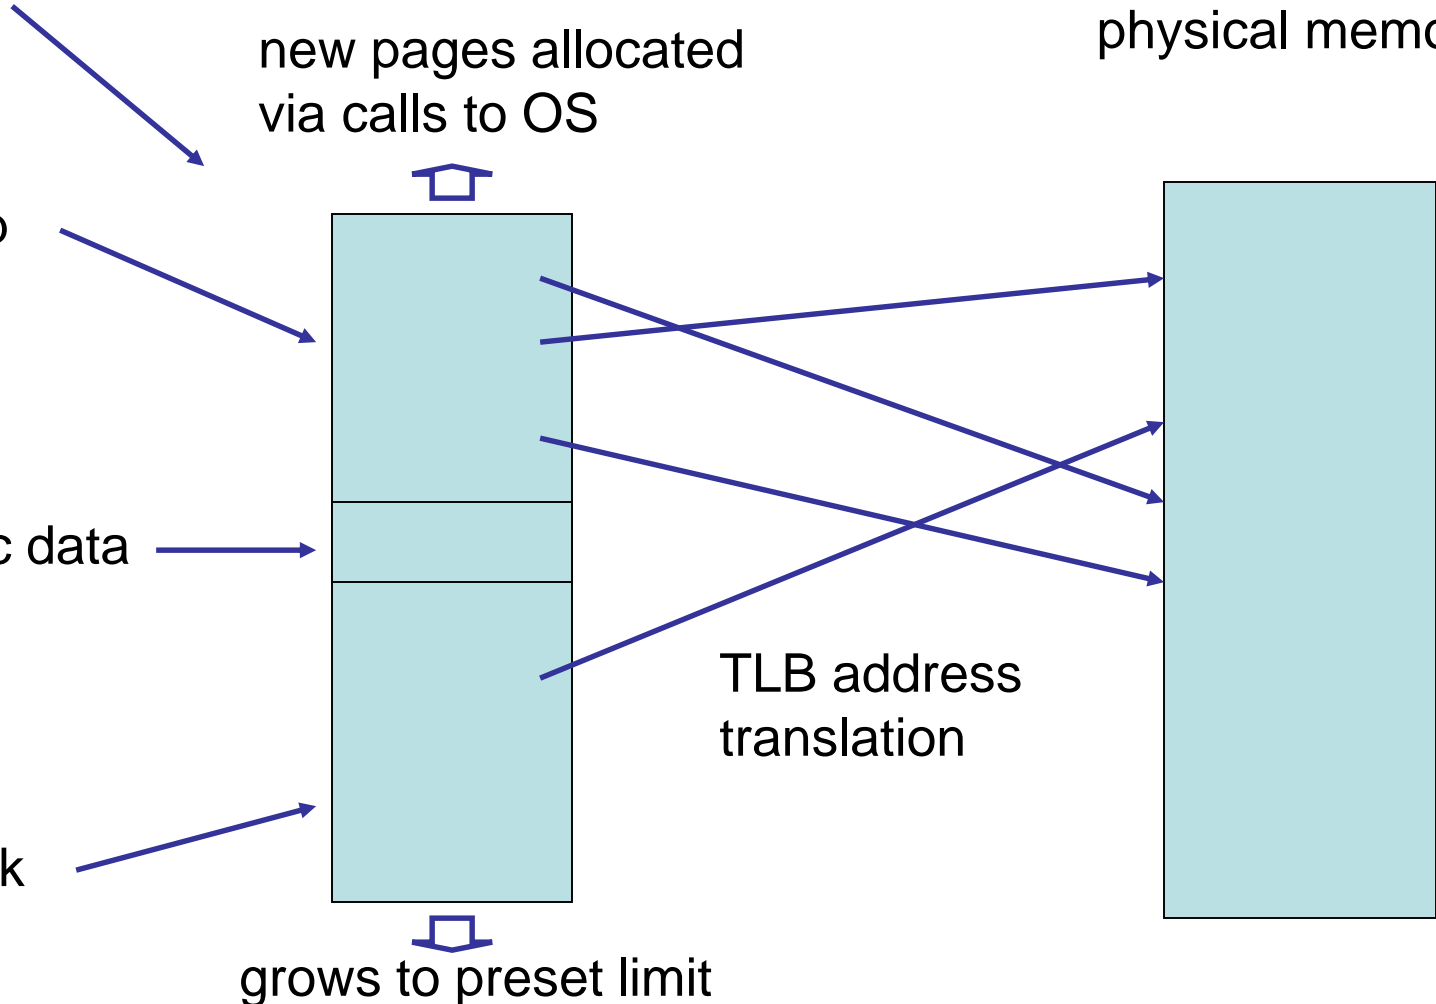
heap

static data

stack

grows to preset limit

TLB address  
translation



# GC

- What is **garbage**?
  - A value is garbage if it **will not be used** in any subsequent computation by the program
- Is it easy to determine which objects are garbage?

# GC

- What is **garbage**?
  - A value is garbage if it **will not be used** in any subsequent computation by the program
- Is it easy to determine which objects are garbage?
  - No. It's undecidable. Eg:  
if long-and-tricky-computation then **use v**  
else **don't use v**

# GC

- Since determining which objects are garbage is tricky, people have come up with many different techniques
  - It's the programmers problem:
    - Explicit allocation/deallocation
  - Reference counting
  - Tracing garbage collection
    - Mark-sweep, copying collection
    - Generational GC


# Explicit MM

- User library manages memory; programmer decides when and where to allocate and deallocate
  - `void* malloc(long n)`
  - `void free(void *addr)`
  - Library calls OS for more pages when necessary
  - Advantage: people are smart
  - Disadvantage: people are dumb and they really don't want to bother with such details if they can avoid it

# Explicit MM

- How does malloc/free work?
  - Blocks of unused memory stored on a **freelist**
  - **malloc**: search free list for usable memory block
  - **free**: put block onto the head of the freelist

freelist






# Explicit MM

- Drawbacks
  - **malloc is not free**: we might have to do a significant search to find a big enough block
  - As program runs, the heap **fragments** leaving many small, unusable pieces

freelist



# Explicit MM

- Solutions:
  - Use multiple free lists, one for each block size
    - Malloc and free become  $O(1)$
    - But can run out of size 4 blocks, even though there are many size 6 blocks or size 2 blocks!
  - Blocks are powers of 2
    - Subdivide blocks to get the right size
    - Adjacent free blocks merged into the next biggest size


# Automatic MM

- Languages with explicit MM are much harder to program than languages with automatic MM
  - Always worrying about **dangling pointers**, **memory leaks**: a huge software engineering burden
  - Impossible to develop a **secure system**, impossible to use these languages in emerging applications involving **mobile code**
  - languages with unsafe, explicit MM will all but **disappear????**

# Automatic MM

- Question: how do we decide which objects are garbage?
  - We conservatively approximate
  - Normal solution: an object is garbage when it becomes unreachable from the roots
    - The roots = registers, stack, global static data
    - If there is no path from the roots to an object, it cannot be used later in the computation so we can safely recycle its memory

# Object Graph



- How should we test reachability?


# Reference Counting

- Keep track of the number of pointers to each object (**the reference count**).
- When the reference count goes to 0, the object is unreachable garbage

# Copying Collection

- Basic idea: use 2 heaps
  - One used by program
  - The other unused until GC time
- GC:
  - Start at the roots & traverse the reachable data
  - Copy reachable data from the active heap (from-space) to the other heap (to-space)
  - Dead objects are left behind in from space
  - Heaps switch roles


# Copying Collection






# Copying GC

- Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Cheny's algorithm for copying collection
  - Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Cheny's algorithm for copying collection
  - Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Cheny's algorithm for copying collection
  - Traverse data breadth first, copying objects from from-space to to-space




# Copying GC

- Cheny's algorithm for copying collection
  - Traverse data breadth first, copying objects from from-space to to-space



# Copying GC

- Traverse data breadth first, copying objects from from-space to to-space





# Copying GC

- Pros
  - Simple & collects cycles
  - Run-time proportional to # live objects
  - Automatic compaction eliminates fragmentation
- Cons
  - Precise type information required
    - Tag bits take extra space; normally use header word
  - Twice as much memory used as program requires
  - Long GC pauses = bad for interactive, real-time apps

# Generational GC

- **Empirical observation:** if an object has been reachable for a long time, it is likely to remain so
- **Empirical observation:** in many languages, most objects died young
- **Conclusion:** we save work by scanning the young objects frequently and the old objects infrequently

# Generational GC

- Assign objects to different generations G0, G1, ...
  - G0 contains young objects, most likely to be garbage
  - G0 scanned more often than G1
  - Common case is two generations (new, tenured)

# Root to G0

- Roots for GC of G0 include all objects in G1 in addition to stack, registers?
- Card marking
- Write Barrier

# Generational GC

- Other issues
  - When do we **promote** objects from young generation to old generation
    - Usually after an object survives a collection, it will be promoted
  - When do we collect the old generation?
    - After several **minor collections**, we do a **major collection**
  - Tuning for efficiency: how quickly do we scan the young generation?

# Generational GC

- Other issues
  - Sometimes different GC algorithms are used for the new and older generations.
    - Why? Because they have different characteristics
  - Copying collection for the new
    - Less than 10% of the new data is usually live
    - Copying collection cost is proportional to the live data
  - Mark-sweep for the old
    - Next topic


# Mark-Sweep

# Mark-sweep

- A two-phase algorithm
  - **Mark phase**: Depth first traversal of object graph from the roots to mark live data
  - **Sweep phase**: iterate over entire heap, adding the unmarked data back onto the free list




# Example



# Example

Mark Phase: mark nodes reachable from roots




Free list




# Example

Mark Phase: mark nodes reachable from roots



# Example

Mark Phase: mark nodes reachable from roots




Free list




# Example

## Sweep Phase: set up sweep pointer; begin sweep




# Example

Sweep Phase: add unmarked blocks to free list




# Example

Sweep Phase



# Example


Sweep Phase: retain & unmark marked blocks





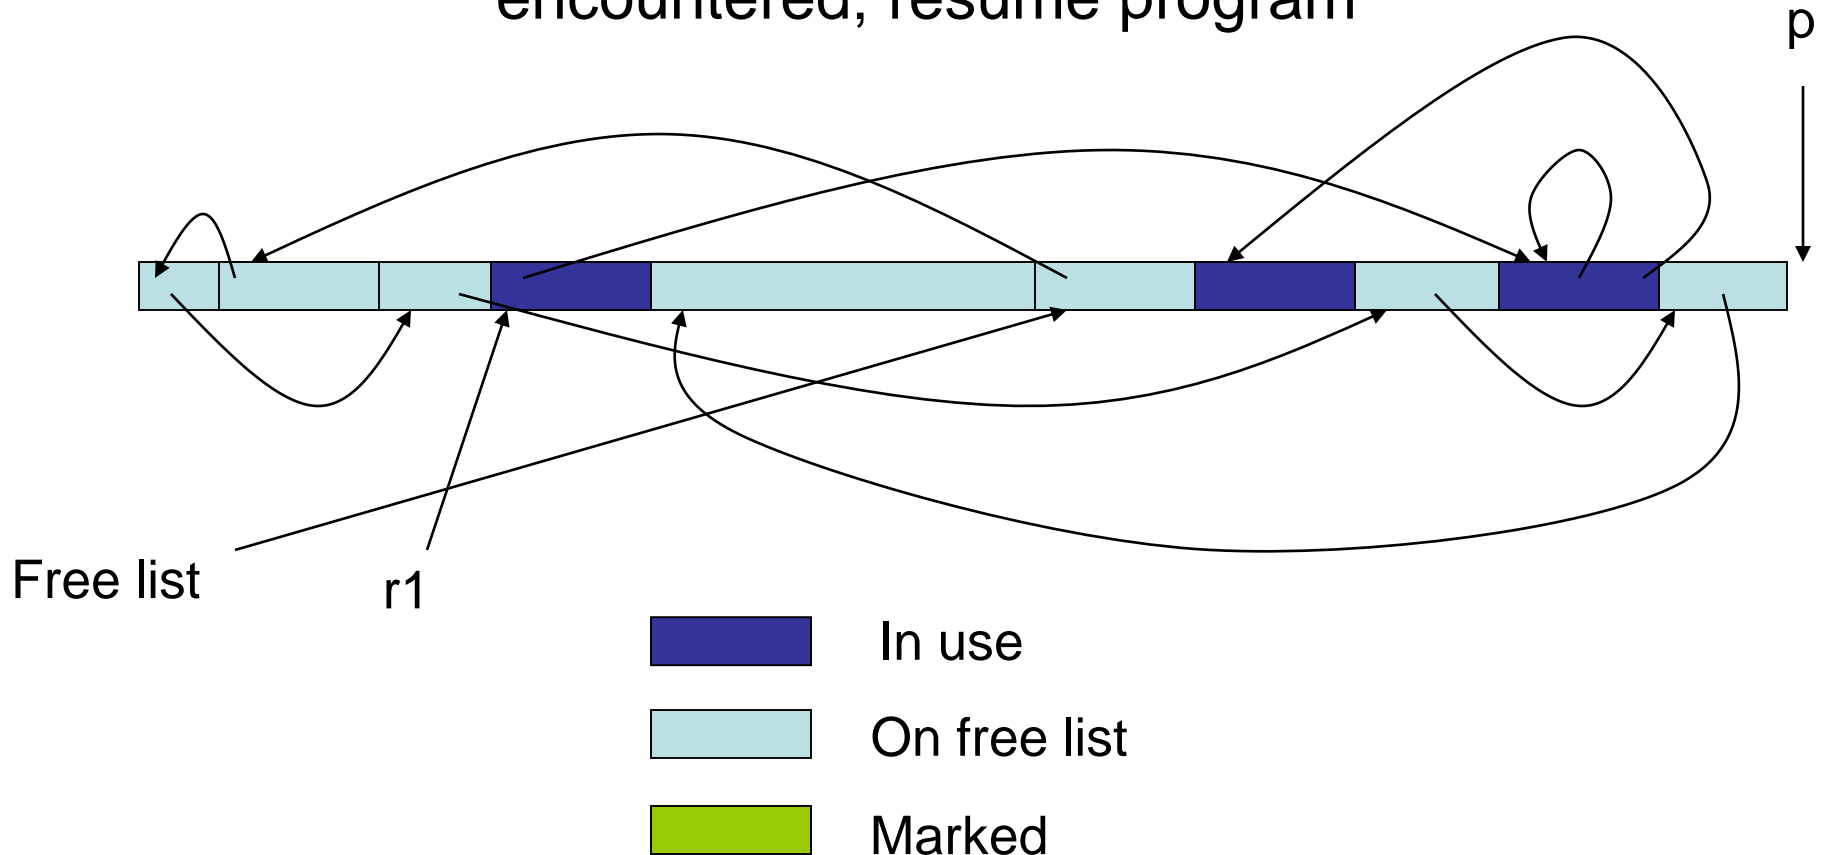
# Example

Sweep Phase



# Example

Sweep Phase: GC complete when heap boundary encountered; resume program



# Discussion

- Relationship between GC cpu usage and size of heap / live set
  - General efficiency (GC cpu / Over-all cpu)
  - Frequency of pauses
  - Length of pause?
    - Copy
    - Mark-sweep

# CMS

- Concurrent Mark Sweep GC
- Stop-the-world copy new generation
- Concurrent, non-compacting old gen
  - Concurrent vs parallel
  - Concurrent marking
    - Lost Object
    - Stop-the-world remark
  - Concurrent sweeping
- Stop-the-world Compacting

# Garbage First (G1)

- New default in JDK 9
- Fine-grained control (target pause time)
  - 2048 regions
  - Eden / Survivor / Tenured
- Stop-the-world copy Young GC
- Concurrent marking
- Incremental compaction Piggybacked on Young GC
  - Do easier ones first (mostly garbage)